

Plan — Domain Product Descriptions: Content Writing Session

Created: 2026-06-16 — Claude.ai Status: Queued — return to this when ready to write

Depends on: `(handoff_domain_product_descriptions_2026-06-16.md)` (Claude Code research, complete)

What this is

Claude Code has done the source research across all three domains and produced a clear gap analysis. The writing work is queued here for a dedicated Claude.ai session. Robert reviews drafts, then a Claude Code spec per domain writes/updates the actual `(docs/)` files.

Not a build spec yet. This becomes three build specs after the writing session.

The three-session output target

One product description per domain, ~500–900 words each, structured:

1. **What it is** — one crisp paragraph
2. **What it does** — the core loop or workflow, user's perspective
3. **Why it's built this way** — design reasoning, not architecture
4. **Current state** — what's live, what's version-tagged, what's next

Destination files (for Claude Code after review):

- `(docs/CURATOR.md)` (new or replace)
 - `(docs/MEIN_DEUTSCH.md)` (tighter replacement for the 6,400-word `(docs/GERMAN.md)`)
 - `(docs/GUILD.md)` (concept-first rewrite of the sprawling mechanics doc)
-

Per-domain brief — what to write toward

Curator

Source anchor: `(docs/curator-research-concept-v2.md)` for philosophy, README for current snapshot, landing blurb for locked voice.

The gap: No single doc says clearly what Curator does, who it's for, and why it's built this way. Concept docs describe a partially-built future state. Current nav `((DAILY · READING`

ROOM · SCANS & DIVES · LEANINGS · ARCHIVE · DESK)) gives the structure but not the story.

What the description needs to land:

- The signal-vs-noise framing: ~700 candidates scored daily, not a feed
- The learned profile: it scores against *your* leanings, not a generic ranking
- The counterpoint design: not an echo chamber by architecture, not by intention
- The two-agent Deeper Dive (Synthesizer + Challenger, CIA Devil's Advocacy)
- Why source architecture matters: trust tiers, age decay, X post hard cap
- In daily use since February 2026 — this is real, not a demo

Locked one-liner to stay consistent with:

~700 candidates scored daily by a reasoning model against your learned profile — counterpoints by design, not an echo chamber. In daily use since February 2026.

Questions to resolve in writing session:

- Is "geopolitics and finance" still the right scope description, or has it broadened?
 - How much of the source architecture is worth surfacing vs. keeping internal?
 - Does the Research Intelligence / Deeper Dive feature warrant its own section or fold into "What it does"?
-

Mein Deutsch

Source anchor: GitHub release (`mein-deutsch-v1.1`) (just published — exact and current). (`docs/GERMAN.md`) for feature depth. Locked landing blurb.

The gap: (`docs/GERMAN.md`) is 6,400 words and predates v1.1 voice/KI-Sitzung. Need a tighter current doc that reflects the cross-workflow pipeline and the Lesen/Gespräche/Schreiben/Wörter/Archiv structure.

What the description needs to land:

- Build your own personas and scenes (not pre-packaged content)
- KI-Personas and real people share one feedback pipeline — the architecture is the differentiator
- The learning loop: practice offline → apply with real people → feed transcripts back → come back better prepared
- Daily-refreshed reading lists tuned to your own interests
- Mistakes become drills automatically — nothing to configure

- Vienna-tested — authenticity signal, not a marketing line

Locked one-liner to stay consistent with:

German practice built around your own life. Build your own KI-Personas and scenes, manage reading lists tuned to your interests, and practice with KI-Personas or real people — transcripts flow through the same feedback pipeline either way. Mistakes become drills. Vienna-tested.

Note: Mein Deutsch product description draft already exists from 2026-06-15 session (docs_mein_deutsch_product_description_2026-06-15.md). Review and tighten against the locked blurb and v1.1 release notes before treating as final. May need only light editing rather than a full rewrite.

Questions to resolve in writing session:

- Does the 2026-06-15 draft hold up against the v1.1 release notes, or does it need structural changes?
 - How much of the tab structure (Lesen/Gespräche/Schreiben/Wörter/Archiv) surfaces in the product description vs. staying in feature docs?
-

Guild

Source anchor: GUILD.md (root) for charter and metaphor. docs/GUILD.md for current state (skim — it's 15k words). Landing blurb for locked voice.

The gap: Existing docs are thorough on mechanics but thin on the *why it's designed this* way angle. Need concept-first: why does an agent system need a "Guild"? What is the Chief of Staff actually doing day to day? What does working with the Guild look like?

What the description needs to land:

- The guild metaphor: craftspeople improving at their work together, not just task execution
- The Chief of Staff role: coordinator with full context, not a domain agent — keeps Robert's real goals and private concerns in view
- Cross-domain by design: Guild sees across Curator, German, Career, and future domains — that's the point
- Current state: Operations agent running under launchd, CoS chat on rvsopenbot, Phase 3 complete
- The "get better at our craft" principle — Guild improves the system itself, not just runs it

Locked one-liner to stay consistent with:

Where the agents and I work together to get better at our craft. From ideas through specs, build, and operations. Guild is intentionally cross-domain. A Chief of Staff role coordinates the agents while keeping my real goals and private concerns in view.

Questions to resolve in writing session:

- Is v0.9 still the right version marker, or has enough shipped to call it v1.0?
- How much of the cabinet model (Chief of Staff, Design-Build, Operations, New Ventures) surfaces publicly vs. stays internal?
- Does Guild need a "who this is for" section — just Robert now, but designed to scale?
- Is the guest access / CoS nudge system far enough along to mention in current state?

Session format — how to run the writing session

Open with: paste this doc + `handoff_domain_product_descriptions_2026-06-16.md`

Work order: Curator first (no existing draft, clearest scope). Mein Deutsch second (existing draft to review, lightest lift). Guild third (most to figure out on the public-facing story).

Per domain:

1. Claude.ai drafts ~500-900 words using source material briefs above
2. Robert reviews inline — line edits or directional notes
3. Lock the draft
4. Claude.ai produces a Claude Code spec to write/update the `docs/` file

One domain per pass — don't try to do all three in one session.

Build specs to produce (after writing session)

Spec	Output file	Depends on
<code>spec_curator_product_description.md</code>	<code>docs/CURATOR.md</code>	Curator draft locked
<code>spec_mein_deutsch_product_description.md</code>	<code>docs/MEIN_DEUTSCH.md</code>	Mein Deutsch

Spec	Output file	Depends on
		draft locked
<code>spec_guild_product_description.md</code>	<code>docs/GUILD.md</code>	Guild draft locked

Each spec includes: the locked draft as content, target file path, DoD (renders correctly in GitHub, consistent with landing blurb, replaces or supplements named existing file), and Commit message.

Files to have ready for the writing session

File	Where
<code>handoff_domain_product_descriptions_2026-06-16.md</code>	This session's uploads
<code>docs_mein_deutsch_product_description_2026-06-15.md</code>	Last session's outputs
This planning doc	This session's outputs
Current landing page screenshot	Take fresh screenshot before session

Planning Doc · 2026-06-16 · Claude.ai Return to this to open the Curator/Guild/Mein Deutsch writing session