

Spec — minimoi.ai Always-On Landing Page (Investigation + Roadmap Entry)

mini-moi · Guild Created: 2026-06-13 — Claude.ai For: Claude Code

Why this matters now

minimoi.ai currently runs via Cloudflare tunnel to Robert's MacBook. The landing page we just polished (What this is / What's running / About) is the front door for job search — LinkedIn and GitHub both point here, and interviewers may check it at any time, including when the Mac is asleep or off (e.g., travel). A portfolio site that's sometimes down undercuts the work just done on it.

This is **investigation + roadmap entry only** — not an implementation spec. The actual fix depends on findings below, and would be a follow-up spec.

Investigation checklist

1. **How is the landing page currently served?** Part of the portal Flask app (port 5001)? A static template? Does it require the backend/DB at all, or is it pure HTML/CSS/text?
 2. **Any dynamic content on the landing page?** E.g., does "~700 candidates scored daily" or similar ever pull live numbers, or is everything hand-written copy (as it is now)? If dynamic, that's a dependency that would need to move too, or be dropped/hardcoded for the static version.
 3. **Current DNS / Cloudflare setup for minimoi.ai** — apex domain routing through the tunnel to the Mac. What would it take to route the root (landing page) to a static host while keeping `/dashboard` (and any other authenticated routes) routed to the tunnel? Options likely include: Cloudflare Pages for the static site + tunnel only for `/dashboard` subpath or subdomain (e.g. `app.minimoi.ai`), or similar.
 4. **Candidate static hosts** — Cloudflare Pages is a natural fit given Cloudflare is already in the stack (tunnel + DNS), but GitHub Pages / Netlify / Vercel are also options. Note tradeoffs if any stand out (e.g., Cloudflare Pages keeping everything in one provider vs. simplest to set up).
 5. **"View source on GitHub" and "LinkedIn" links** — confirm these are plain external links with no backend dependency (should be trivial, but confirm as part of the audit).
-

Working hypothesis (for Claude Code to confirm or correct)

Split the landing page into a static page (no Flask, no backend) hosted separately — up 24/7 regardless of the Mac's state. "Dashboard →" continues to link to the live Flask app via the Cloudflare tunnel, which is up whenever the Mac is on. Interviewers/recruiters always see the front door; the Dashboard being occasionally unavailable just means a logged-in feature is offline, not the whole site.

If investigation finds this is harder than expected (e.g., landing page is tightly coupled to the Flask app's templating/data), note that and propose alternatives — including the previously-deferred Mac Mini migration (v1.2 on the roadmap) as a different way to solve "always on" if the static-split approach isn't clean.

Definition of Done

- ☐ Investigation checklist above completed, findings written up (in this file or a follow-up doc — Claude Code's call)
- ☐ Recommendation given: static-split (with rough plan), Mac Mini migration, or other — with reasoning
- ☐ Entry added to (docs/ROADMAP.md) under "Next" (or "Platform-Portfolio" if that's still the better-fitting section), along the lines of:

```
> **minimoi.ai always-on landing page** — Public landing page
> currently depends on Cloudflare tunnel + MacBook uptime. For job
> search, needs to be reliably available independent of Mac state.
> See `spec_landing_page_uptime_2026-06-13.md` for investigation
> findings and recommended approach.
```

- ☐ If recommendation is clear and small, a follow-up implementation spec can be written in the same pass — Robert's call whether to do that now or treat as separate spec-ready item
-

Spec · Landing Page Uptime · 2026-06-13