

mini-moi

mini-moi is a personal AI agent platform. It carries context across reading and research, language learning, and the work of building and running the system itself. It remembers my interests, what I have already investigated, how my thinking develops, what I decide to follow up on, and what I got wrong.

The same pattern applies beyond one person. Persistent context and feedback can help teams, organizations, and their systems learn and improve over time. This repository is the personal version.

mini-moi has been in daily use since February 2026, on AWS since June 2026. This repository documents the system as it works today: architecture, operations, roadmap, and ongoing development, with current capability kept distinct from planned work.

Current state (August 2026)

Domain	Current role
Curator	Supports daily reading and sustained research in geopolitics and finance. Each morning, it distills roughly 700 candidates from international RSS feeds, X bookmarks, and theme-driven web search into a curated shortlist of 20. Selected articles can become scans, research threads, or deeper dives. Reactions, comments, and recorded leanings shape later selections. A serendipity pool guards against a closed loop, while AI synthesis surfaces patterns across the accumulated record.
Mein Deutsch	Originally inspired by an upcoming trip to Vienna and a lack of opportunities to speak German, this domain uses AI personas in everyday scenarios for unscripted conversations adapted to the learner's ability. The intent is immersion, including the tension of following and speaking in a new language. It was road-tested during the trip itself, when unexpected communication challenges could be turned into new scenes and relived later. Reading and writing reinforce the learning loop and preserve locally grounded context. The learner's activity then shapes future sessions.
Meu Português	Portuguese was chosen to fill a family need: maintaining a second language used at home across varying levels of proficiency. AI personas, scenes, and reading content are Brazilian-focused. The domain does not put users into fixed proficiency levels. The intent is still immersion. A more advanced learner can choose a more demanding scenario and speak or write in a more complex way. A novice can use simpler scenarios and rely more heavily on targeted drills or traditional vocabulary practice. Usability enhancements were added across the site, including hover translation, tips in English, and the option to speak throughout.
Guild	Guild was created to bring discipline, visibility, and learning to the development of the platform. Build, Operate, and Improve are its three working areas. Build currently maintains a Kanban view of work in progress, an inventory of specifications, and an evolving roadmap. With more than one coding agent now in use and operational concerns becoming more important, the planned Master Craftsman will coordinate build quality, code review, domain standards, and consistency between specifications and implementation.
Chief of Staff	Chief of Staff is intended to become my personal strategic thinking and cross-domain working partner. It will carry context across mini-moi, notice connections, support decisions, and help direct follow-through. Its scope will also extend to learning and professional priorities beyond the current domains, using the platform's shared tools and capabilities. The beta supports discussion, investigation, and refinement of possible actions but cannot alter the platform itself. OpenClaw is the current bounded agent, running independently on AWS behind a stable interface so it can be replaced without redesigning the domain. Once proven, the same approach will support Master Craftsman , a separate build-focused agent in Guild.

Production runs on AWS EC2, with a Mac development standby. The deployment consists of nine application/runtime containers plus PostgreSQL, with host-level Cloudflare Tunnel, nginx, and cron. Production currently uses a full immutable-image, health-gated deployment. A reviewed but not yet shipped CI/CD change adds document-only and domain-scoped releases so documentation can publish without container restarts and one domain can deploy without interrupting the others; shared or ambiguous changes continue to select the full path. Error tracking and scheduled health checks monitor the running system.

The August 16 milestone is deliberately called a **beta**, not 1.0. It proves a real agentic loop—conversation, bounded research, model routing, verified note mutation, and realtime voice—while longer-term conversation retention, privacy controls, natural capture intents, cross-domain reach, and cost governance remain explicit follow-up work.


The maintained reference documents are:

- [ARCHITECTURE.md](#) - system design and rationale
- [OPERATIONS.md](#) - how the production system runs
- [ROADMAP.md](#) - current plans and the record of what shipped

The [product tour](#) shows the current production interface across all five domains and follows one Curator topic from the daily briefing into long-form research.

Two product loops

Curator and the language domains are the clearest expressions of the product. One turns daily reading into a continuing research practice. The other turns an unscripted conversation into material for the next session.




The colors separate the two domains; the double-braced nodes mark the places where a model interprets activity rather than simply storing or displaying it.


Curator: from reading to research

A selected article can remain a quick read, contribute feedback to later selection, or become the starting point for a longer investigation.

Reading and feedback



Research




Feedback operates at two timescales. Immediate actions - likes, dislikes, saves, and comments - help shape later article scoring. Research threads, deeper dives, and recorded leanings preserve how an issue develops across sources and sessions. The second loop is less mature than article selection, so the roadmap and architecture describe its current limits.

Language practice: conversation and return

German began with an upcoming trip to Vienna and a lack of opportunities to speak before going. The domain addresses three practical gaps. First, I need the experience of getting stuck and recovering in a real exchange. Second, I want to read engaging, current material without the effort becoming burdensome; contextual hints, translation, audio, and shorter excerpts help me keep moving. Third, I want reasons to write about something I actually find interesting, with correction and explanation for the mistakes I make. Gespräche, Lesen, and Schreiben are complementary responses to those needs rather than separate exercises.

The persona opens the exchange. The conversation is free-form, and difficulty inside the conversation becomes evidence for later practice rather than a reason to replace it with a scripted lesson. The approach was road-tested in Vienna, where unexpected communication challenges could become new scenes to relive later.




German came first and remains more developed in the backend, including the use of reviewed conversation context in later persona prompts. Portuguese followed to meet a family need: maintaining a second language used at home across varying levels of proficiency. Each person begins with the same experience and keeps a separate learning record. The domain remains level-agnostic: advanced learners can choose more demanding scenes and language, while novices can rely more on simpler scenarios, hover translation, vocabulary practice, and targeted drills.

The two domains are intended to converge into the same solution from interface through stored data and adaptive memory. A third language, new to the whole family, will test that design directly: the language and content should change; the solution should not.

Lesen: reading, writing, and reuse

Lesen uses current German-language articles from subjects with an existing reason to read: **Alltag** (everyday life), **Kultur** (culture), **Nachrichten** (news), and **Wien** (Vienna and local life). A selected article opens with a short excerpt drawn from the beginning of the source. The full article remains available on its original site.



Translation help and other contextual hints are available without replacing the reading. They keep unfamiliar language from turning an interesting article into a burdensome decoding exercise. A highlighted word or phrase can be translated and added to the Wörter library with its source sentence and article title. The excerpt can also be expanded or read aloud.

Writing remains tied to the subject being read. A response can be typed or dictated, and the correction step returns natural German, an English translation, and short explanations of the changes. The corrected version can be accepted or revised before saving. The archive retains the article, the original response, and the corrected response; a drill record is saved with the same work.

Lesen supplies the slower layers that real-time conversation cannot: close reading, contextual translation, time to formulate a response, correction, and deliberate review. Those layers develop recognition and recall in service of better real conversation. Speaking exposes gaps; reading and writing provide material and repetition; later conversation tests whether the language can be retrieved.

Continuity while the system changes

mini-moi is intended to keep its history even as the technology changes. Five years from now, the useful part should be the knowledge gathered, lessons learned, language practice, decisions, and some continuity of character from daily interaction - not any particular model, memory system, interface, or hosting arrangement. Provider memory can be used when it helps, while application and local records keep that history from depending on one provider.

Domains give the work practical boundaries. Curator, the language domains, Guild, and Chief of Staff each have a defined purpose and set of guardrails. This reduces confusion about what a task is trying to accomplish and limits the effect of mistakes. The boundaries can change: new interests may become domains, while older areas may need less attention.

Curator and German were useful starting points because they supported real interests and repeated use while also providing a way to learn AI by building with it. Parts of the continuity loop already work, including Curator's article selection, German's reviewed conversation memory, and Lesen's saved vocabulary and corrected writing. Longer-term measurement and fuller coordination are still being built. The aim is not to preserve today's components, but to let mini-moi adapt without losing the history that makes it a more useful collaborator. [ROADMAP.md](#) records that boundary as it changes.

Design approach

A few choices explain most of the architecture.

Keep learned state outside the model. Personal context lives in files owned by the application rather than inside one model provider. The files are structured to match a later database schema, and migration paths have already been used in the Guild and research layers. Infrastructure is added when the volume or behavior requires it.

Use different models for different work. Model control rests on three considerations:

- **Cost.** Model costs became significant, then dropped substantially when different tasks were routed to different models, including a local option with no per-call charge.
- **Quality for the task and evaluation.** Models differ in conversation, translation, analysis, review, and coding, and those differences change over time. The ability to switch models also supports direct A/B evaluation on the same task.
- **Vendor independence and continuity.** Personal context remains above the model layer, so changing providers does not reset the history or require a redesign of the personalization logic. Provider-managed memory can supplement that record without becoming its only home.

Applying model choice consistently at every call site remains ongoing work; a July audit found several places where implementation still lagged the intended configuration.

The production history reflects these tradeoffs. The first version ran on a local model, and a local model remains the last resort in German's translation fallback chain. Curator scoring moved to cloud models after the cost and quality tradeoff was measured. The path back to local execution remains part of the design and is scheduled for end-to-end verification on the current EC2 setup.



Use AI agents as contributors, with human direction. The director defines the intent, decides what enters the system, and remains responsible for the result. Claude Code, Claude.ai, OpenClaw, Grok, and OpenAI Codex contribute planning, design alternatives, implementation, testing, and review. Separately, the Chief of Staff beta introduces OpenClaw as the first bounded agent operating inside the product. The particular agents and models will change; the separation of roles and the platform-owned boundaries are more durable. Important changes receive review from more than one agent before they are treated as complete.




Operate while the system is evolving. Keeping daily use dependable is important and not trivial, especially while development continues. Drift can appear between domains, between code and documentation, and between assumed safeguards and the protection that actually exists during an incident. Operations is therefore a continuing process: observe production, test assumptions, learn from failures, reconcile code and documentation, and close gaps in safeguards as they are found.

Product tour

These production captures from July 21, 2026 show the five domain environments and several workflows behind their landing pages. The descriptions reflect the current August 16 state, including the newer Chief of Staff beta. Each image opens at full size.



Five domains



Curator	Mein Deutsch
	
Daily reading, research, working views, and the long-term archive.	German conversation, reading, writing, and vocabulary in a place-based learning environment.

Meu Português	Guild	Chief of Staff
		
Family Portuguese practice with separate learning records, level-agnostic immersion, and voice throughout.	Build, Operate, and Improve make platform development and operations visible.	A production beta for conversation, bounded investigation, and platform-owned notes; the image predates the agent release.

Curator: daily reading becomes research



The daily briefing is the beginning of the Curator loop, not the finished product. A ranked article can be read quickly, used to provide feedback, scanned for its argument and evidence, or turned into a larger investigation.

Daily briefing	Article scan
	
The morning selection ranks articles and keeps Like, Pass, and Save close to the reading decision.	The scan connects one article to the question that made it interesting, then extracts its argument, counterpoint, evidence, and next questions.

Research controls	Deeper-dive result
	
Bibliography items can seed further research; the article can also become an immediate deeper dive or a multi-session thread.	The completed synthesis states what the research did and did not establish, preserves uncertainty, and includes a separate challenge pass.

Language practice

The language interfaces use photography, personas, and real subjects to give practice a setting and a reason. Correction and review remain attached to the learner's own writing or conversation rather than becoming isolated exercises.

German writing and correction	Portuguese conversation setup
	
A fresh demonstration sentence is corrected and explained without exposing a saved learning session.	A persona, scenario, model, and voice workflow prepare an unscripted Portuguese conversation.

Earlier screenshots remain in [docs/screenshots/](#) as part of the project's history. Approved sets are grouped by date so a later visual refresh can be added without overwriting this one.

More detail

- [docs/README.md](#) is the guide to working references and historical context, including the April 2026 platform vision.
 - [docs/journal/](#) is the tracked home for the public build narrative. Its publishing conventions are present; dated sprint notes, plans, and records of how the project changed are still being populated.
 - [docs/releases/](#) contains release notes.
 - [ARCHITECTURE.md](#), [OPERATIONS.md](#), and [ROADMAP.md](#) are the maintained core documents.
-

Status: In daily use since February 2026; running on AWS since June 2026.

Author: Robert van Stedum