

mini-moi

Download the formatted PDF

mini-moi is a personal AI agent platform. It carries context across reading and research, language learning, and the work of building and running the system itself. It remembers my interests, what I have already investigated, how my thinking develops, and what I decide to follow up on.

The same pattern applies beyond one person. Persistent context and feedback that compounds instead of resetting is just as useful to a team or an organization that wants its systems to learn from repeated work rather than start over each time. This repository is the personal version.

mini-moi has been in daily use since February 2026, on AWS since June 2026. This repository documents the system as it works today: architecture, operations, roadmap, and ongoing development, with current capability kept distinct from planned work.

Current state (July 2026)

Domain	Current role
Curator	Supports daily reading and longer research in geopolitics and finance. Each morning it reduces roughly 700 candidates from international RSS feeds, X bookmarks, and theme-driven web search to 20 items on the web and 10 in Telegram. Reactions, saved items, comments, and investigations provide context for later selection. A separate serendipity pool helps prevent the learned profile from becoming a closed loop. Curator can also produce article scans, research threads, deeper-dive briefs, AI observations, and a weekly synthesis.
Mein Deutsch	Combines simulated German immersion with interest-led reading and writing. Gespräche provides unscripted voice conversations with AI personas and user-chosen scenes. Lesen presents current German-language articles across everyday life, culture, news, and Vienna; each selection includes a short source excerpt and a link to the full original. Words and phrases can be translated and saved with context. Article-related notes can be typed or dictated, corrected and explained by AI, then retained with their original and corrected forms. Vocabulary, writing records, conversation reviews, and the archive provide material for repeated use across the different forms of practice.
Meu Português	Followed German as an intentional copy-and-adapt extension for a mixed English-Portuguese household. Both parents use both languages, but each has a different stronger language; the children have different levels of fluency, exposure, and formal practice. The family shares a desire to maintain and develop those abilities over time. The first release deliberately emphasized a polished frontend, multi-user separation, and per-user data boundaries so real users could evaluate it quickly. The backend is now converging with German, especially in adaptive memory.
Guild	Gives me one place to follow ongoing planning, development, and operations across the platform. That visibility becomes more important as multiple coding agents contribute build work. Its Build, Operate, and Improve sections remain in use while the domain is refocused after Chief of Staff moved out. The planned Master Craftsman will coordinate build quality, code review, domain standards, and consistency between specifications and implementation inside Guild, without mixing that role with Chief of Staff.
Chief of Staff	Released as v0.9 after becoming a standalone domain in July 2026. It provides chat with a defined voice, its own working memory and agenda, scheduled watch loops, and health checks. It remains in use while its remaining dependencies on Guild are removed and its access across the other domains is expanded. The separation gives both Chief of Staff and Guild clearer boundaries for further development.

Production runs on AWS EC2, with a Mac development standby. The deployment consists of eight containers plus host-level nginx and cron. CI/CD normally moves a pushed change to production in about five minutes. Error tracking and scheduled health checks monitor the running system.


The maintained reference documents are:

- [ARCHITECTURE.md](#) - system design and rationale
- [OPERATIONS.md](#) - how the production system runs
- [ROADMAP.md](#) - current plans and the record of what shipped

The [product tour](#) shows the current production interface across all five domains and follows one Curator topic from the daily briefing into long-form research.

Two product loops

Curator and the language domains are the clearest expressions of the product. One turns daily reading into a continuing research practice. The other turns an unscripted conversation into material for the next session.




The colors separate the two domains; the double-braced nodes mark the places where a model interprets activity rather than simply storing or displaying it.


Curator: from reading to research

A selected article can remain a quick read, contribute feedback to later selection, or become the starting point for a longer investigation.

Reading and feedback



Research




Feedback operates at two timescales. Immediate actions - likes, dislikes, saves, and comments - help shape later article scoring. Research threads, deeper dives, and recorded leanings preserve how an issue develops across sources and sessions. The second loop is less mature than article selection, so the roadmap and architecture describe its current limits.

Language practice: conversation and return

For my German, the domain addresses three practical gaps. First, I need more chances to speak, including the experience of getting stuck and recovering in a real exchange. Second, I want to read engaging, current material without the effort becoming burdensome; contextual hints, translation, audio, and shorter excerpts help me keep moving. Third, I want reasons to write about something I actually find interesting, with correction and explanation for the mistakes I make. Gespräche, Lesen, and Schreiben are complementary responses to those needs rather than separate exercises.

The persona opens the exchange. The conversation is free-form, and difficulty inside the conversation becomes evidence for later practice rather than a reason to replace it with a scripted lesson.




German came first and remains more developed in the backend, including the use of reviewed conversation context in later persona prompts. Portuguese followed as an intentional copy-and-adapt extension for family use in a mixed English-Portuguese household. It moved quickly and deliberately focused first on the frontend, multi-user separation, and per-user data boundaries so it could enter real family use and extend the language pattern beyond German.

The two domains are intended to converge into the same solution from interface through stored data and adaptive memory. A third language, new to the whole family, will test that design directly: the language and content should change; the solution should not.

Lesen: reading, writing, and reuse

Lesen uses current German-language articles from subjects with an existing reason to read: **Alltag** (everyday life), **Kultur** (culture), **Nachrichten** (news), and **Wien** (Vienna and local life). A selected article opens with a short excerpt drawn from the beginning of the source. The full article remains available on its original site.



Translation help and other contextual hints are available without replacing the reading. They keep unfamiliar language from turning an interesting article into a burdensome decoding exercise. A highlighted word or phrase can be translated and added to the Wörter library with its source sentence and article title. The excerpt can also be expanded or read aloud.

Writing remains tied to the subject being read. A response can be typed or dictated, and the correction step returns natural German, an English translation, and short explanations of the changes. The corrected version can be accepted or revised before saving. The archive retains the article, the original response, and the corrected response; a drill record is saved with the same work.

Lesen supplies the slower layers that real-time conversation cannot: close reading, contextual translation, time to formulate a response, correction, and deliberate review. Those layers develop recognition and recall in service of better real conversation. Speaking exposes gaps; reading and writing provide material and repetition; later conversation tests whether the language can be retrieved.

Continuity while the system changes

mini-moi is intended to keep its history even as the technology changes. Five years from now, the useful part should be the knowledge gathered, lessons learned, language practice, decisions, and some continuity of character from daily interaction - not any particular model, memory system, interface, or hosting arrangement. Provider memory can be used when it helps, while application and local records keep that history from depending on one provider.

Domains give the work practical boundaries. Curator, the language domains, Guild, and Chief of Staff each have a defined purpose and set of guardrails. This reduces confusion about what a task is trying to accomplish and limits the effect of mistakes. The boundaries can change: new interests may become domains, while older areas may need less attention.

Curator and German were useful starting points because they supported real interests and repeated use while also providing a way to learn AI by building with it. Parts of the continuity loop already work, including Curator's article selection, German's reviewed conversation memory, and Lesen's saved vocabulary and corrected writing. Longer-term measurement and fuller coordination are still being built. The aim is not to preserve today's components, but to let mini-moi adapt without losing the history that makes it a more useful collaborator. [ROADMAP.md](#) records that boundary as it changes.

Design approach

A few choices explain most of the architecture.

Keep learned state outside the model. Personal context lives in files owned by the application rather than inside one model provider. The files are structured to match a later database schema, and migration paths have already been used in the Guild and research layers. Infrastructure is added when the volume or behavior requires it.

Use different models for different work. Model control rests on three considerations:

- **Cost.** Model costs became significant, then dropped substantially when different tasks were routed to different models, including a local option with no per-call charge.
- **Quality for the task and evaluation.** Models differ in conversation, translation, analysis, review, and coding, and those differences change over time. The ability to switch models also supports direct A/B evaluation on the same task.
- **Vendor independence and continuity.** Personal context remains above the model layer, so changing providers does not reset the history or require a redesign of the personalization logic. Provider-managed memory can supplement that record without becoming its only home.

Applying model choice consistently at every call site remains ongoing work; a July audit found several places where implementation still lagged the intended configuration.

The production history reflects these tradeoffs. The first version ran on a local model, and a local model remains the last resort in German's translation fallback chain. Curator scoring moved to cloud models after the cost and quality tradeoff was measured. The path back to local execution remains part of the design and is scheduled for end-to-end verification on the current EC2 setup.

Use AI agents as contributors, with human direction. The director defines the intent, decides what enters the system, and remains responsible for the result. Different agents contribute design alternatives, implementation, testing, and review. The current set includes Claude Code, Claude.ai, OpenClaw, Grok, and OpenAI Codex. The particular models will change; the separation of roles is more durable. Important changes receive review from more than one agent before they are treated as complete.

Operate while the system is evolving. Keeping daily use dependable is important and not trivial, especially while development continues. Drift can appear between domains, between code and documentation, and between assumed safeguards and the protection that actually exists during an incident. Operations is therefore a continuing process: observe production, test assumptions, learn from failures, reconcile code and documentation, and close gaps in safeguards as they are found.

Product tour

These production captures from July 21, 2026 show the five domains as they are currently presented and several of the workflows behind their landing pages. Each image opens at full size.



Five domains



Curator	Mein Deutsch
Daily reading, research, working views, and the long-term archive.	German conversation, reading, writing, and vocabulary in a place-based learning environment.

Meu Português	Guild	Chief of Staff
A family-use Portuguese version with per-user boundaries and a polished front end.	The build, operations, and improvement workspace.	A standalone coordination domain, still visibly marked as in development.

Curator: daily reading becomes research



The daily briefing is the beginning of the Curator loop, not the finished product. A ranked article can be read quickly, used to provide feedback, scanned for its argument and evidence, or turned into a larger investigation.

Daily briefing	Article scan
	
The morning selection ranks articles and keeps Like, Pass, and Save close to the reading decision.	The scan connects one article to the question that made it interesting, then extracts its argument, counterpoint, evidence, and next questions.

Research controls	Deeper-dive result
	
Bibliography items can seed further research; the article can also become an immediate deeper dive or a multi-session thread.	The completed synthesis states what the research did and did not establish, preserves uncertainty, and includes a separate challenge pass.

Language practice

The language interfaces use photography, personas, and real subjects to give practice a setting and a reason. Correction and review remain attached to the learner's own writing or conversation rather than becoming isolated exercises.

German writing and correction	Portuguese conversation setup
 A fresh demonstration sentence is corrected and explained without exposing a saved learning session.	 A persona, scenario, model, and voice workflow prepare an unscripted Portuguese conversation.

Earlier screenshots remain in [docs/screenshots/](#) as part of the project's history. Approved sets are grouped by date so a later visual refresh can be added without overwriting this one.

More detail

- [docs/README.md](#) is the guide to working references and historical context, including the April 2026 platform vision.
- [docs/journal/](#) is the tracked home for the public build narrative. Its publishing conventions are present; dated sprint notes, plans, and records of how the project changed are still being populated.
- [docs/releases/](#) contains release notes.
- [ARCHITECTURE.md](#), [OPERATIONS.md](#), and [ROADMAP.md](#) are the maintained core documents.

Status: In daily use since February 2026; running on AWS since June 2026.
Author: Robert van Stedum